

Definition

Project Overview

Nearly half of the world depends on seafood for their main source of protein. In the Western and Central Pacific, where 60% of the world's tuna is caught, illegal, unreported, and unregulated fishing practices are threatening marine ecosystems, global seafood supplies and local livelihoods. [The Nature Conservancy](#) is working with local, regional and global partners to preserve this fishery for the future.

Currently, the Conservancy is looking to the future by using cameras to dramatically scale the monitoring of fishing activities to fill critical science and compliance monitoring data gaps. Although these electronic monitoring systems work well and are ready for wider deployment, the amount of raw data produced is cumbersome and expensive to process manually.

The Conservancy started a competition on the [Kaggle](#) community to develop algorithms to automatically detect and classify species of tunas, sharks and more that fishing boats catch, which will accelerate the video review process. Faster review and more reliable data will enable countries to reallocate human capital to management and enforcement activities which will have a positive impact on conservation and our planet.

Problem Statement

We have to detect which species of fish appeared on the boat based on images that has been captured from the boat cameras of various angles.

Our goal is to predict the likelihood of the fish species in each picture. Eight target categories are available in the dataset: Albacore tuna, Bigeye tuna, Yellowfin tuna, Mahi Mahi, Opah, Sharks, Other (meaning that there are fish present but not in the above categories) and No Fish (meaning that no fish is in the picture). Each image has only one fish category, expect that there are sometimes very small fish in the image that are used as a bait.

Metrics

The Evaluation metrics used for this classification problem is multi-class logarithmic loss.

$$\log\text{-loss} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^M y_{ij} \log(p_{ij})$$

Where N is the number of images in the test set, M is the number of class labels of image, log is the natural logarithm, y_{ij} is 1 if the observation i belongs to class j and 0 otherwise, p_{ij} is the predicted probability that observation i belongs to class j .

The probabilities for a given image is not required to sum to one as they are scaled prior to being scored, each row is divided by the row sum. In order to avoid the extremes of the log function, predicted probabilities are replaced by $\max(\min(p, 1 - 10^{-15}), 10^{-15})$.

Analysis

Data Exploration

For this competition we, have around 4000 images for training dataset and 1000 images for testing dataset, captured from boat cameras of various angles.

Eight Target Categories are available in the dataset

1. Albacore tuna
2. Bigeye tuna
3. Yellowfin tuna
4. Mahi Mahi
5. Opah
6. Sharks
7. Other (meaning that there are fish present but not in the above categories)
8. No Fish (meaning that no fish is in the picture)



ALB: Albacore tuna (*Thunnus alalunga*)



BET: Bigeye tuna (*Thunnus obesus*)



DOL: Dolphinfish, Mahi Mahi (*Coryphaena hippurus*)



LAG: Opah, Moonfish (*Lampris guttatus*)



SHARK: Various: Silky, Shortfin Mako



YFT: Yellowfin tuna (*Thunnus albacares*)

Fish images are not to scale with one another

The dataset was compiled by [The Nature Conservancy](#) in partnership with [Satlink](#), [Archipelago Marine Research](#), the [Pacific Community](#), the [Solomon Islands Ministry of Fisheries and Marine Resources](#), the [Australia Fisheries Management Authority](#), and the governments of [New Caledonia](#) and [Palau](#).

Exploratory Visualization

1. In the provided Dataset, there are lot of images with significant variations in the color intensity as the images were taken from different angles of the camera and different time of a day. And there are few thousand images available for the training and this can be a difficult task to classify these images correctly.

The Color images have 3 channels RGB i.e. R-Red, G-Green, B-Blue. The below images provide the color distribution of the image. Red, Green and Blue line represents the respective channels of an image.

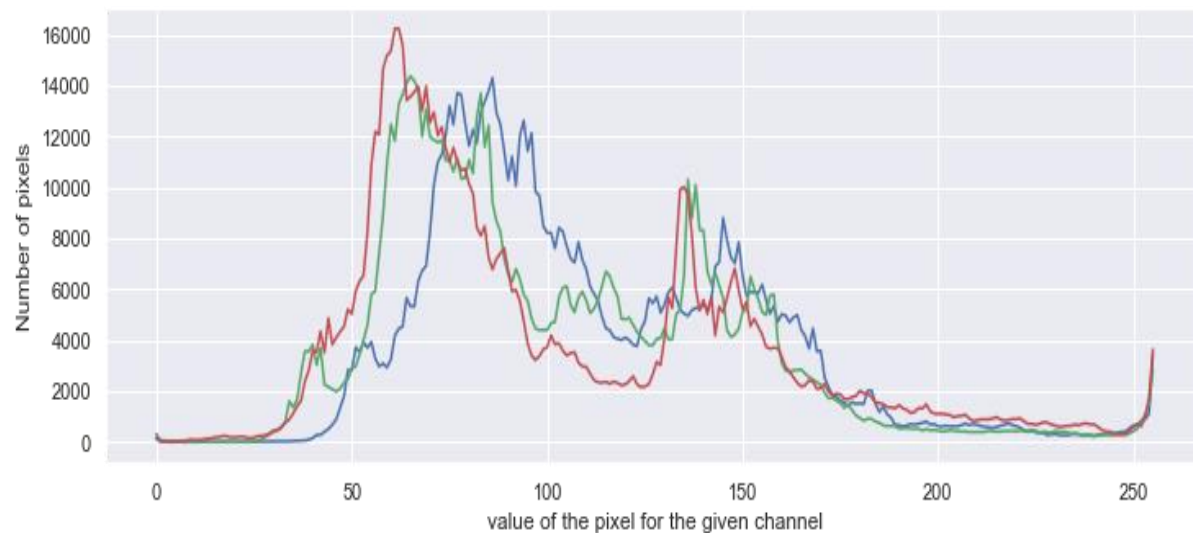



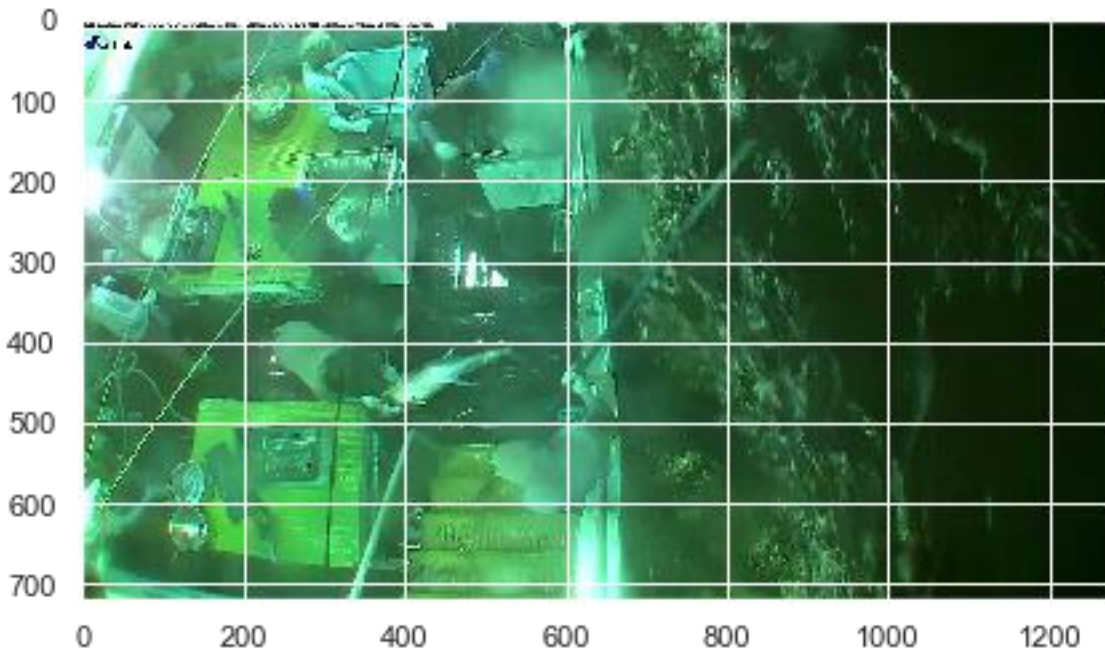
Fig: Demonstrated RGB distribution of below image



**Fig: Random Image taken from training dataset
Image Size (675 X 1200)**



Fig: RGB distribution of below image



**Fig: Random image taken from training dataset
Image Size (720 X 1280)**

2. It can be observed from the above pictures that fish is present in only a small part of the entire image and the rest of the part is useless which we don't need to classify. It is seen from the pictures that fish is present in only a small part of the image and the rest is useless stuff which we do not need to classify the fish.
3. The training data provided for this competition was not distributed uniformly among the different classes of the fishes.




Fig: Visualization of number of Fishes in each class of training dataset

Number of fishes in each class of the training dataset as shown in the above figure.

ALB	1719
BET	200
DOL	117
LAG	67
NoF	465
OTHER	299
SHARK	176
YFT	734

Table: Number of Fishes in each class of Training Dataset

It may be problematic, as we have a large number of samples from 'ALB' class and very few samples from 'LAG' class the model may find it difficult to make a correct prediction for the later class and predict an image as 'ALB' more often.

Algorithms and Techniques

1. Deep Learning - Deep learning also known as deep structured learning or hierarchical learning is part of a broader family of machine learning methods based on learning data representations, as opposed to task-specific algorithms.

2. Convolutional neural network - In machine learning, a convolutional neural network

(CNN or ConvNet) is a class of deep, feed-forward artificial neural networks (ANN) that has successfully been applied to analyzing visual imagery. A CNN consists of an input and an output layer with multiple hidden layers. The hidden layers are either convolutional, pooling or fully connected layers. We give CNN an input and it learns by itself that what features it has to detect. We won't specify the initial values of features or what kind of patterns it has to detect.

Various Layers involved, as follows:

- **Convolutional layer** - Also referred to as Conv. layer, it forms the basis of the CNN and performs the core operations of training and consequently firing the neurons of the network. It performs the convolutional operation over the input.

- **Pooling layer** - Pooling layer reduces the spatial dimensions (Width x Height) of the input Volume for the next Convolutional Layer. It does not affect the depth dimension of the Volume.

- **Fully connected layer** - The fully connected or Dense layer is configured exactly the way its name implies. It is fully connected with the output of the previous layer. Fully connected layers are typically used in the last stages of the CNN to connected to the output layer and construct the desired number of outputs.

- **Dropout layer** - Dropout is a regularization technique for reducing overfitting in neural networks by preventing complex co-adaptations on training data. It is a very efficient way of performing model averaging with neural networks. The term "dropout" refers to dropping out units (both hidden and visible) in a neural network.

- **Flatten** - Flattens the output of the convolutional layers to feed into the Dense layers.

3. Activation Functions - In CNN, the activation function of a node defines the output of that node given an input or set of inputs. Some activation functions are, as follows:

- **Softmax** - Softmax function squashes the output of each unit to be between 0 and 1, just like a sigmoid function. It also divides each output such that the total sum of the outputs is equal to 1.

- **ReLU** - A ReLU (or rectified linear unit) has output 0 if the input is less than 0, and raw output otherwise. If the input is greater than 0, the output is equal to the input.

4. Transfer Learning - In transfer learning, we take the learned understanding and pass it to a new deep learning model. We take a pre-trained neural network and adapt it to a new neural network with different dataset.

For this problem we use VGG-19 neural network.

VGG-19: It is a 19-layer network used by the VGG team in the ILSVRC-2014 competition.

Benchmark model

For this competition, a CNN is made from the scratch, it is also used as a benchmark model for measuring the performance of transfer learning approach.




Fig: Benchmark Model Architecture

Benchmark model Architecture

- **Conv2D** - Two Conv2D layers were used, the first layer has 16 and the second layer has 32 filters. Kernel size and strides were set to 2 and 1 respectively and both of the layers use a 'relu' activation function.
- **MaxPooling2D** - One MaxPooling2D layer with pool size and strides both equal to 2.
- **GlobalAveragePooling2D** - One GlobalAveragePooling2D layer was used before the Dense layer.
- **Dropout** - Two dropout layers are used both having the value of 0.3.

- **Dense** - One fully connected dense layer with 8 nodes (equal to the number of classes of fish) and 'softmax' activation function were used at the end.

After defining the architecture for benchmark model, it was trained on the training set with the validation split of 0.2 and the best weights were saved during the training process. After training, predictions were made on test set and it achieved a multi-class log loss score of **1.51541** on public leaderboard and **1.75980** on private leaderboard.

Methodology

Data Preprocessing

I have used Keras CNNs, it requires a 4D array (4D tensors) as input with shape nb_samples, rows, columns and channels; where 'nb_samples' is the total number of samples, 'rows' is number of rows, 'columns' is number of columns and 'channels' is channels for each image. The input images in the data set differed in the dimension so we need to resize them and convert them into a 4D tensor. A 'path_to_tensor' function in the script takes a string valued file path to a color image as input and returns a 4D tensors. The function loads the image and resizes it to (224 X 224) pixels. Then the image is converted to an array which is then resized to a 4D tensors.

Each tensor corresponds to an image and have a shape (1,244,244,3) where nb_samples = 1, rows = 224, columns = 224, channels = 3 (color images have 3 channels i.e. Red, Green and Blue color).

Implementation

1. An initial **benchmark CNN model** architecture was defined and it was trained on the pre-processed training data using validation split of 0.2. The training data was divided into batches of 20 images and the best weights were saved using check-pointer. Then the labels of the test images were predicted and saved into a .csv file for submission. The benchmark model achieved a Multi-class log loss score of **1.51541** on public leaderboard and **1.75980** on private leaderboard.

2. In Model 2 the **transfer learning** was used to create a CNN using **VGG-19** bottleneck features. VGG-19 architecture without the last fully connected layer was used to extract the features from the images. The extracted features were fed to a small fully connected model and predictions were made. It takes around 5-6 hours for extracting the features from the images and making the predictions. The model did not perform well and get a score of **1.64077** on public leaderboard and **4.04953** on private leaderboard.

3. In Model 3, I had reduced the number of layers in VGG-19 by removing 11 layers from the end as I thought that it may be a case where Model 2 was giving more emphasis to the other details than the fish in the image. When I tried to extract the features of the images this time my machine ran out of memory and resulted into **Memory Error**.

4. After first three try, I decided to create another model from scratch. This time I used a combination of Conv2D, MaxPooling, Dropout, GlobalMaxPooling and Dense layers and was able to get a multi-class log score of **1.50787** on public leaderboard and **1.76167** on private leaderboard.

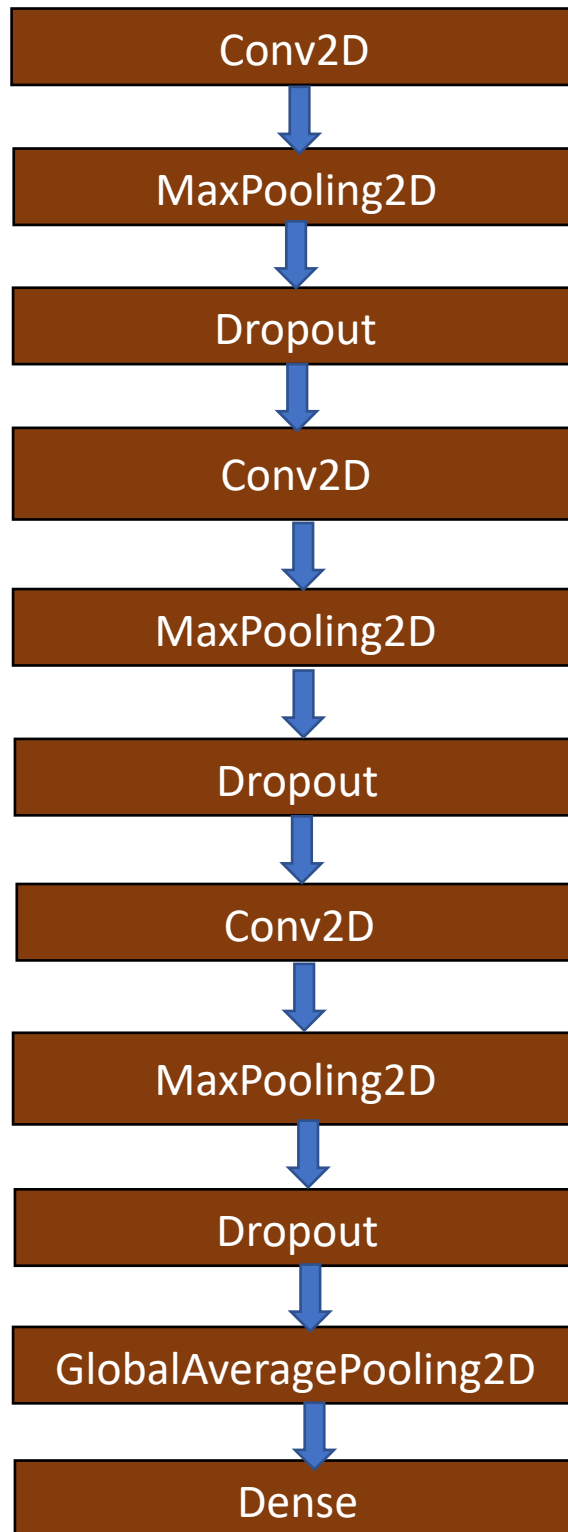


Fig: Architecture of Final Model

Refinement

- Model4 - Final model shows a significant improvement when compared with the benchmark model. I use 3 Conv layers with increasing number of filters from 16 in first, 32 in second to 64 filters in the last conv layer. After each Conv and MaxPooling, a Dropout layer with probability 0.2 was inserted. This model was able to achieve a multi-class log loss score of **1.50787** on public leaderboard and **1.76167** on private leaderboard.
- Model5 - Then, I double the number of filters in each Conv layer and increase the value of first and second dropout layer so that it takes less time for training. I train this model for only 5 epochs and able to get a score of **1.50961** on public leaderboard and **1.77940** on private leaderboard.
- Model6 - Seeing this improvement, I further try to improve my model, this time changing the value of all the dropout layers back to 0.2 and using 'rmsprop' optimizer instead of 'adam'. I train this model for 10 epochs and get a multi-class log loss score of **1.51393** on public leaderboard and **1.70159** on private leaderboard. **This model would get in top 5% in this Kaggle competition with 1.70159 score on private leaderboard.**



Results

Model Evaluation and Validation

The final model outperforms the benchmarks and other models used in this project. It takes nearly half of the time taken by the model which uses transfer learning approach (with extracted VGG-19 features), and make better predictions.

The parameters of this model were carefully adjusted so that it can make better predictions. It can be reflected by the multi-class log loss score achieved by the model.

For this competition, Kaggle provides a very large testing set with 13153 images in it. All the images are captured from different boat cameras at different angles and at different times of the day. The dimensions and the quality of the images differed significantly.



Predictions are made on the testing set and then submitted on the Kaggle competition. The score provided by Kaggle for the submission was considered as the norm for measuring the robustness of the model.

Justification

The final model performs better than the benchmark and other models as well. Where

benchmark model gets a multiclass log loss score of **1.75980** on private leaderboard, the final model made a significant improvement of approximately 0.05 and get a score of **1.70159**.

The model does not expect any image of a particular size or taken from a particular angle, it can automatically resize the image and extract the features from it to make its predictions.

It was able to get a decent score when we submit the predictions made by the model on the provided test data set, given that the test dataset has a large number of images of different sizes, taken from different angles at different times of day, and of different quality.

Conclusion

Free-Form Visualization

The data provided for this project are images taken from the boat camera at different times of day and from different angles. The images differ in their sizes and quality, some images are captured in very poor lighting conditions and some pictures are very bright and clear. The training images are not distributed uniformly among the different classes, 'ALB' class has nearly half of the images of the whole training set.




Fig: Shows number of training samples in each class

As we have a large number of samples from 'ALB' class and very few samples from 'LAG' class the model may find it difficult to make a correct prediction for the later class and predict an image as 'ALB' more often.

Reflection

In the beginning I build a benchmark model and then tried to improve the result using transfer learning technique with VGG-19 model.

Initially, when I proposed this project I thought that transfer learning was really easy to implement and always outperform the other methods, but as I proceed through my project I got to know that I was wrong. It is not easy to apply Transfer learning approach to any solution until

you have a clear understanding of it and a lot of hands on practice. It is computationally very expensive to train big models on CPUs and requires GPUs which can run several processes parallelly.

Then I shifted back from transfer learning to create my own model architecture and made improvements to it by fine tuning the parameters.

The final model performs better than my benchmark model as well as better than Kaggle's benchmark for this competition and now I have a better understanding of some concepts about CNN and deep learning and definitely, I will look forward doing some more projects to learn more about them.

Improvement

This was my first experience using CNN and I found that they require a lot more computing power than any traditional approaches and we really need to take care of the memory consumptions, this is one of the important take away from this project as I ran out of memory for several times.

Data augmentation could also be used to increase the number of training samples and the model could be trained for more numbers of epochs.

I still think that transfer learning approach can outperform my final model. Maybe other architectures like VGG-16 and Resnet-50 with carefully adjusted parameters by experimentation can produce a better multi-class log loss score.

References

<https://www.kaggle.com/c/the-nature-conservancy-fisheries-monitoring>

<https://www.conserveca.org/>

<https://www.nature.org/>

<https://satlink.es/en/>

<https://www.archipelago.ca/>

<https://www.spc.int/>

<https://fisheries.gov.sb/>

<https://www.afma.gov.au/>

<https://gouv.nc/>

<https://palaugov.pw/executive-branch/ministries/natural-resources/>

https://en.wikipedia.org/wiki/Transfer_learning

<http://cs231n.github.io/convolutional-networks/>

[https://en.wikipedia.org/wiki/Dropout_\(neural_networks\)](https://en.wikipedia.org/wiki/Dropout_(neural_networks))

<http://cs231n.github.io/transfer-learning/>

<https://www.analyticsvidhya.com/blog/2016/10/tutorial-optimizing-neural-networks-usingkeras-with-image-recognition-case-study/>

<https://keras.io/>

<https://in.udacity.com/>